


project 14: Markov chains

In this project, we are concerned with homogeneous discrete time Markov chains over discrete finite state spaces $S = \{s_1, s_2, \dots, s_m\}$. That is, we are concerned with stochastic processes where

$$p(X_t = s_j \mid X_{t-1} = s_i) = p(s_j \mid s_i) = P_{ji}.$$

In the lecture we looked at examples such as this one (which models the behavior of “guarding the treasure in room B ”)



$$P = \begin{bmatrix} 0.25 & 0.10 & 0.25 \\ 0.50 & 0.80 & 0.50 \\ 0.25 & 0.10 & 0.25 \end{bmatrix}$$

Recall that the matrix of state transition probabilities of a homogeneous DTMC is column stochastic. In other words, if $|S| = m$, then $P \in \mathbb{R}^{m \times m}$ and, importantly

$$P_{ji} \geq 0 \quad \wedge \quad \sum_j P_{ji} = 1.$$

Here is a puzzling statement: a Markov chain evolves over a state space and has states itself. But the states a Markov is in are not the states over which it evolves. In fact, states of a Markov chain are stochastic vectors $\pi \in \mathbb{R}^m$ where

$$\pi_i \geq 0 \quad \wedge \quad \sum_i \pi_i = 1.$$

The significance of these state vectors is as follows: if we let $\pi[t]$ denote the state vector at time t , we have

$$p(X_t = s_i) = \pi_i[t].$$

In other words, $\pi[t]$ is a probability distribution over the set of states S and $\pi_i[t]$ indicates how likely the stochastic process is in state s_i after t time steps.

State vectors at time t are computed as follows: given an initial state distribution $\pi[0]$, we have

$$\begin{aligned}\pi[1] &= P\pi[0] \\ \pi[2] &= P\pi[1] = PP\pi[0] \\ \pi[3] &= P\pi[2] = PPP\pi[0] \\ &\vdots \\ \pi[t] &= P\pi[t-1] = \underbrace{P \cdot P \cdots P}_{t \text{ times}} \pi[0] = P^t \pi[0]\end{aligned}$$


If this process converges for $t \rightarrow \infty$, it reaches a so called stationary distribution π where

$$\pi = P\pi.$$

In other words, if a Markov chain has a stationary distribution π , then π is an eigenvector of P and the corresponding eigenvalue is 1.

task 14.1: evolution of a Markov process

Consider again the Markov model for “guarding the treasure in room B ”



$$P_2 = \begin{bmatrix} 0.25 & 0.10 & 0.25 \\ 0.50 & 0.80 & 0.50 \\ 0.25 & 0.10 & 0.25 \end{bmatrix}$$

Given P_2 as above, compute $\pi[t]$ for


$$t \in \{1, 2, 4, 8, 16\}$$

and

$$\pi[0] \in \left\{ \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \right\}$$

What do you observe? Interpret your results. Answer the question: “what is the most likely state s_i for the stochastic process to be in after t steps?”

Next, consider our Markov model for “patrolling the map”



$$P_1 = \begin{bmatrix} 0.30 & 0.20 & 0.50 \\ 0.50 & 0.30 & 0.20 \\ 0.20 & 0.50 & 0.30 \end{bmatrix}$$

and, given P_1 , compute the state distributions $\pi[t]$ just as above. What do you observe in this case? Interpret your results.

Finally, compute the spectral decompositions

$$\Lambda_1 = U_1 P_1 U_1^\top$$

$$\Lambda_2 = U_2 P_2 U_2^\top$$

and have a look at the resulting eigenvector matrices U_1, U_2 . What do you observe?

task 14.2: sampling a Markov process

Here is another puzzling statement: there is a difference between the most likely state of stochastic process to be in at time t and the state it is actually in at time t .

To see what this means, create a state sequence according to transition matrix P_2 .

In the supplementary material for lecture 19, we discussed that, given an appropriate numpy array P_2 , this can be accomplished as follows:

```
import numpy as np
import numpy.random as rnd

states = ['A', 'B', 'C']
indices = range(len(states))

state2index = dict(zip(states, indices))
index2state = dict(zip(indices, states))

def generateStateSequence(X0, P, tau):
    sseq = [X0]

    iold = state2index[X0]

    for t in range(tau):
        inew = rnd.choice(indices, p=P[:,iold])
        sseq.append(index2state[inew])
        iold = inew

    return sseq
```

To generate and print a sequence of length 10 that starts with state B, use

```
>>> sequence = generateStateSequence('B', P2, 9)
>>> print (''.join(sequence))
```

In fact, generate 10.000 sequences of length 10 and have a look at their respective last element `sequence[-1]`. What do you observe? Given this sample of sequences, what is the empirically likeliest last element? How likely is it?